

This is a test PDF for layout-aware parsing using PyMuPDF and pdfplumber.

Multiple lines of text to test chunking and section preservation.

A second paragraph demonstrating PyMuPDF layout-aware extraction at scale.

Recursive character chunker with sliding-window overlap walks separators in priority order: paragraph, then sentence, then word. Each emitted chunk preserves the source page number so the LLM retains spatial context.

Phase C.1 introduces PyMuPDF + pdfplumber with pypdf as a graceful fallback.

The dispatch in loader.py adds a `prefer_v2` kwarg defaulting to `True`, matching the INC-005 zero-runtime-change constraint from the 32-commit DDD effort.

Total: roughly 700 characters across 10 lines, which clears the 500 chars/page threshold set by the inspector, so the routing decision is "text".